

Session 7

Incorporating covariates into LCM

Matt Denwood

2022-06-10

Recap

Models for diagnostic test evaluation require:

- At least 2 tests
- At least 2 populations, but preferably 3 or more
- Quite a lot of data

Recap

Models for diagnostic test evaluation require:

- At least 2 tests
- At least 2 populations, but preferably 3 or more
- Quite a lot of data

Fitting the models is technically quite straightforward

The real difficulty lies in the interpretation

- What exactly is the latent class?

Incorporating covariates for prevalence

Modelling variation in infection probability

- Individuals may be at higher/lower risk of being infected due to known characteristics e.g.:
 - Age
 - Sex
 - History
 - Presence of co-infections
 - Whatever

Modelling variation in infection probability

- Individuals may be at higher/lower risk of being infected due to known characteristics e.g.:
 - Age
 - Sex
 - History
 - Presence of co-infections
 - Whatever
- There are three ways to deal with this:
 1. Ignore it
 2. Group “populations” by these characteristics
 3. Embed a (preferably simple!) generalised linear model within your LCM

Logistic regression in JAGS

```
model{

  for(i in 1:N){
    Observation[i] ~ dbern(prob[i])
    logit(prob[i]) <- intercept + beta1[Category[i]] + beta2*Covariate[i]
  }

  intercept ~ dnorm(0, 0.01)
  beta1[1] <- 0
  for(c in 2:NC){
    beta1[c] ~ dnorm(0, 0.01)
  }
  beta2 ~ dnorm(0, 0.01)

  #data# N, Observation, NC, Category, Covariate
  #monitor# intercept, beta1, beta2
  #inits# intercept, beta1, beta2
}
```

```

model{

  for(i in 1:N){
    Observation[i] ~ dbern(obs_prob[i])
    obs_prob[i] <- prob[i]*se + (1-prob[i])*(1-sp)
    logit(prob[i]) <- intercept + beta1[Category[i]] + beta2*Covariate[i]
  }

  se ~ dbeta(148.43, 16.49)T(1-sp, )
  sp ~ dbeta(240.03, 12.63)

  intercept ~ dnorm(0, 0.01)
  beta1[1] <- 0
  for(c in 2:NC){
    beta1[c] ~ dnorm(0, 0.01)
  }
  beta2 ~ dnorm(0, 0.01)

  #data# N, Observation, NC, Category, Covariate
  #monitor# intercept, beta1, beta2, se, sp
  #inits# intercept, beta1, beta2, se, sp
}

```



```

model{

  for(i in 1:N){
    Observation[i] ~ dbern(obs_prob[i])
    obs_prob[i] <- prob[i]*se + (1-prob[i])*(1-sp)
    logit(prob[i]) <- intercept + beta1[Category[i]] + beta2*Covariate[i]
  }

  #data# se, sp

  intercept ~ dnorm(0, 0.01)
  beta1[1] <- 0
  for(c in 2:NC){
    beta1[c] ~ dnorm(0, 0.01)
  }
  beta2 ~ dnorm(0, 0.01)

  #data# N, Observation, NC, Category, Covariate
  #monitor# intercept, beta1, beta2
  #inits# intercept, beta1, beta2
}

```

```

model{

  for(i in 1:N){
    Observation[i] ~ dbern(obs_prob[i])
    obs_prob[i] <- prob[i]*se[Test[i]] + (1-prob[i])*(1-sp[Test[i]])
    logit(prob[i]) <- intercept + beta1[Category[i]] + beta2*Covariate[i]
  }

  #data# se, sp

  intercept ~ dnorm(0, 0.01)
  beta1[1] <- 0
  for(c in 2:NC){
    beta1[c] ~ dnorm(0, 0.01)
  }
  beta2 ~ dnorm(0, 0.01)

  #data# N, Observation, NC, Category, Covariate, Test
  #monitor# intercept, beta1, beta2
  #inits# intercept, beta1, beta2
}

```

```

model{

  for(i in 1:N){
    Observations[i,1:4] ~ dmulti(obs_probs[i,1:4], 1)

    obs_probs[i,1] <- (prob[i] * ((1-se[1])*(1-se[2]))) + ((1-prob[i]) * ((sp[1])*(sp[2])))
    obs_probs[i,2] <- (prob[i] * ((se[1])*(1-se[2]))) + ((1-prob[i]) * ((1-sp[1])*(sp[2])))
    obs_probs[i,3] <- (prob[i] * ((1-se[1])*(se[2]))) + ((1-prob[i]) * ((sp[1])*(1-sp[2])))
    obs_probs[i,4] <- (prob[i] * ((se[1])*(se[2]))) + ((1-prob[i]) * ((1-sp[1])*(1-sp[2])))

    logit(prob[i]) <- intercept + beta1[Category[i]] + beta2*Covariate[i]
  }

  #snip#

}

```

```

model{

  for(i in 1:G){
    Observations[i,1:4] ~ dmulti(obs_probs[i,1:4], Total[i])

    obs_probs[i,1] <- (prob[i] * ((1-se[1])*(1-se[2]))) + ((1-prob[i]) * ((sp[1])*(sp[2])))
    obs_probs[i,2] <- (prob[i] * ((se[1])*(1-se[2]))) + ((1-prob[i]) * ((1-sp[1])*(sp[2])))
    obs_probs[i,3] <- (prob[i] * ((1-se[1])*(se[2]))) + ((1-prob[i]) * ((sp[1])*(1-sp[2])))
    obs_probs[i,4] <- (prob[i] * ((se[1])*(se[2]))) + ((1-prob[i]) * ((1-sp[1])*(1-sp[2])))

    logit(prob[i]) <- intercept + beta1[Category[i]] + beta2*RoundedCovariate[i]
  }

  #snip#

}

```

Embedding a LR within a LCM

- Hierarchical modelling gives us the freedom to include “models within models” as needed for our specific application, but we usually have to write these ourselves!
- Remember that blocking at group level is much more efficient than looping through all individuals as fewer likelihood calculations are required

Embedding a LR within a LCM

- Hierarchical modelling gives us the freedom to include “models within models” as needed for our specific application, but we usually have to write these ourselves!
- Remember that blocking at group level is much more efficient than looping through all individuals as fewer likelihood calculations are required
- Autocorrelation may be problematic - if so try to use different contrast schemes eg:

```
sex_effect ~ dnorm(0, 0.01)
beta1[1] <- -sex_effect/2
beta1[2] <- sex_effect/2
```

- Random effects are kind of like fixed effects:

```
#snip#
  logit(prob[i]) <- intercept + beta1[Category[i]] + beta3[Group[i]]
#snip#

for(r in 1:NR){
  beta3[r] ~ dnorm(0, tau)
}
tau ~ dgamma(0.01, 0.01)

#inits# tau
#monitor# tau, beta3
```

Generating code for a LR

You can use template.jags as inspiration:

```
template.jags(weight ~ group, family="gaussian", data=data, file="linear_model.txt")  
## Your model template was created at "linear_model.txt" - it is highly advisable to examine  
↪ the model syntax to be sure it is as intended  
## You can then run the model using run.jags("linear_model.txt")  
results <- run.jags("linear_model.txt")  
## Loading required namespace: rjags  
## module glm loaded  
## module dic loaded
```



```
#####
#####
#### JAGS model file written by runjags version 2.2.2-1 on 2022-06-09 22:18:03
#####
#####
```

```
### Model template as follows - ensure this is syntactically correct before running the
↪ model!
```

```
model{
```

```
# In the BUGS/JAGS language we must use an explicit for loop:
for(i in 1:N){
  # These lines describe the response distribution and linear model terms:
  weight[i] ~ dnorm(regression_fitted[i], regression_precision)
  regression_residual[i] <- weight[i] - regression_fitted[i]
  regression_fitted[i] <- intercept + group_effect[group[i]]
}
```

```
# These lines give the prior distributions for the parameters to be estimated:
regression_precision ~ dgamma(0.001, 0.001)
intercept ~ dnorm(0, 10^-6)
group_effect[1] <- 0      # Factor level "Ctl"
group_effect[2] ~ dnorm(0, 10^-6)    # Factor level "Trt"
resid.sum.sq <- sum(regression_residual^2)
}
```

Supported features:

- Gaussian, binomial, Poisson, negative binomial, ZIB, ZIP, ZINB
- Random intercepts
- Automatic centering of continuous variables

We can also add (currently manually):

- Random slopes
- Spline terms
- Interval censoring

Grouping populations

- This is much easier than writing a GLM within our LCM...
- We can also use the `template_huiwalter` function!
 - See Otero-Abad 2017 for a simple example
- If you have a lot of populations you could use a simple random effect:

```
# prev[p] ~ dbeta(1, 1)
logit(prev[p]) <- intercept + raneff[i]
raneff[i] ~ dnorm(0, tau)
```

Grouping populations

- This is much easier than writing a GLM within our LCM...
- We can also use the `template_huiwalter` function!
 - See Otero-Abad 2017 for a simple example
- If you have a lot of populations you could use a simple random effect:

```
# prev[p] ~ dbeta(1, 1)
logit(prev[p]) <- intercept + raneff[i]
raneff[i] ~ dnorm(0, tau)
```

- Be careful that Se/Sp is still consistent across populations!

Do nothing?

What is the goal of your analysis?

- Estimating risk factors for disease?
- Estimating true prevalence?
- Estimating Se/Sp?

Do nothing?

What is the goal of your analysis?

- Estimating risk factors for disease?
- Estimating true prevalence?
- Estimating Se/Sp?

Inclusion of risk factors for disease is NOT necessary to estimate Se/Sp!

If you are interested in risk factors for disease (rather than the se/sp directly) then I would probably use a simpler model with fixed se/sp (+/- multiple imputation)

Covariates for Sensitivity / Specificity

What if diagnostic tests are not consistent across populations?

This time we can't just ignore it!

Solutions:

- Remove that population (and clearly state this in the paper..!)
- Allow the relevant parameter to vary between populations
- Use a (very simple) GLM on the relevant parameter(s)

What if diagnostic tests are not consistent across populations?

This time we can't just ignore it!

Solutions:

- Remove that population (and clearly state this in the paper..!)
- Allow the relevant parameter to vary between populations
- Use a (very simple) GLM on the relevant parameter(s)

But now we are no longer technically within the Hui-Walter framework. . .

Varying Se/Sp between populations

```
Tally[1:4,p] ~ dmulti(prob[1:4,p], N[p])
```

```
prob[1,p] <- prev[p] * ((1-se[1,SeGp[p]])*(1-se[2,SeGp[p]])) + (1-prev[p]) *  
↪ (sp[1])*(sp[2])
```

```
prob[2,p] <- prev[p] * (se[1,SeGp[p]]*(1-se[2,SeGp[p]])) + (1-prev[p]) *  
↪ (1-sp[1])*sp[2]
```

```
prob[3,p] <- prev[p] * ((1-se[1,SeGp[p]])*(se[2,SeGp[p]])) + (1-prev[p]) *  
↪ (sp[1])*(1-sp[2])
```

```
prob[4,p] <- prev[p] * (se[1,SeGp[p]]*se[2,SeGp[p]]) + (1-prev[p]) *  
↪ (1-sp[1])*(1-sp[2])
```

```
prev[p] ~ dbeta(1,1)
```

Varying Se/Sp between populations

```
Tally[1:4,p] ~ dmulti(prob[1:4,p], N[p])
```

```
prob[1,p] <- prev[p] * ((1-se[1,SeGp[p]])*(1-se[2,SeGp[p]])) + (1-prev[p]) *  
↪ (sp[1])*(sp[2])
```

```
prob[2,p] <- prev[p] * (se[1,SeGp[p]]*(1-se[2,SeGp[p]])) + (1-prev[p]) *  
↪ (1-sp[1])*sp[2]
```

```
prob[3,p] <- prev[p] * ((1-se[1,SeGp[p]])*(se[2,SeGp[p]])) + (1-prev[p]) *  
↪ (sp[1])*(1-sp[2])
```

```
prob[4,p] <- prev[p] * (se[1,SeGp[p]]*se[2,SeGp[p]]) + (1-prev[p]) *  
↪ (1-sp[1])*(1-sp[2])
```

```
prev[p] ~ dbeta(1,1)
```

See also: Stærk-Østergaard 2022

Embedded GLM for Se/Sp

```
for(p in 1:Populations){  
  Tally[1:4,p] ~ dmulti(prob[1:4,p], N[p])  
  
  # Probability of observing test -/-  
  prob[1,p] <- prev[p] * ((1-se[1,p])*(1-se[2,p])) + (1-prev[p]) * (sp[1])*(sp[2])  
  
  #snip#  
  
  logit(se[1,p]) <- se_intercept[1] + se1_beta[Se1Category[p]]  
  logit(se[2,p]) <- se_intercept[2] + se2_beta[Se2Category[p]]  
  
}  
  
# NB: tweak contrasts for convergence!  
se_beta1[1] <- -se_eff/2  
se_beta1[2] <- se_eff/2  
se_eff ~ dnorm(0, 0.01)
```

Embedded GLM for Se/Sp

```
for(p in 1:Populations){  
  Tally[1:4,p] ~ dmulti(prob[1:4,p], N[p])  
  
  # Probability of observing test -/-  
  prob[1,p] <- prev[p] * ((1-se[1,p])*(1-se[2,p])) + (1-prev[p]) * (sp[1])*(sp[2])  
  
  #snip#  
  
  logit(se[1,p]) <- se_intercept[1] + se1_beta[Se1Category[p]]  
  logit(se[2,p]) <- se_intercept[2] + se2_beta[Se2Category[p]]  
  
}  
  
# NB: tweak contrasts for convergence!  
se_beta1[1] <- -se_eff/2  
se_beta1[2] <- se_eff/2  
se_eff ~ dnorm(0, 0.01)
```

See also: Martinez 2008 (although they wrote the samplers themselves!)

Embedded GLM for Infection AND Se/Sp

```
for(p in 1:Populations){
  Tally[1:4,p] ~ dmulti(prob[1:4,p], N[p])

  # Probability of observing test -/-
  prob[1,p] <- prob[p] * ((1-se[1,p])*(1-se[2,p])) + (1-prob[p]) * (sp[1])*(sp[2])

  #snip#

  logit(se[1,p]) <- se_intercept[1] + se1_beta[Se1Category[p]]
  logit(se[2,p]) <- se_intercept[2] + se2_beta[Se2Category[p]]

  logit(prob[i]) <- intercept + beta1[Category[i]]
}

#snip#
se_eff ~ dnorm(0, 0.01)
pr_eff ~ dnorm(0, 0.01)
```

Embedded GLM for Infection AND Se/Sp

```
for(p in 1:Populations){
  Tally[1:4,p] ~ dmulti(prob[1:4,p], N[p])

  # Probability of observing test -/-
  prob[1,p] <- prob[p] * ((1-se[1,p])*(1-se[2,p])) + (1-prob[p]) * (sp[1])*(sp[2])

  #snip#

  logit(se[1,p]) <- se_intercept[1] + se1_beta[Se1Category[p]]
  logit(se[2,p]) <- se_intercept[2] + se2_beta[Se2Category[p]]

  logit(prob[i]) <- intercept + beta1[Category[i]]

}

#snip#
se_eff ~ dnorm(0, 0.01)
pr_eff ~ dnorm(0, 0.01)
```

But what happens if “Category” is confounded with “Se1Category” ??

General points

- Inconsistent se/sp may happen in e.g. laboratory vs field settings, blood vs milk samples, or due to biological effects such as age-acquired immunity
- Theoretically it is possible to incorporate this into the model, but. . .
 - It is not frequently done so the models are less well developed/understood
 - Great care is needed, and any additions should be based on known biological processes
 - If all populations have their own se/sp then the model collapses!

General points

- Inconsistent se/sp may happen in e.g. laboratory vs field settings, blood vs milk samples, or due to biological effects such as age-acquired immunity
- Theoretically it is possible to incorporate this into the model, but. . .
 - It is not frequently done so the models are less well developed/understood
 - Great care is needed, and any additions should be based on known biological processes
 - If all populations have their own se/sp then the model collapses!
- Be VERY careful when prevalence and se/sp have the same covariate
 - It may work if balancing populations by these covariates, and only including them as se/sp covariates?

Practical session 7

Points to consider

1. What is the optimal number of populations?
2. What happens to identifiability when you deviate from the standard Hui-Walter model?

Summary

- Adding populations (or equivalently, covariates on prevalence) adds parameters but may add information
 - But it is not always worthwhile!
- Using covariates on sensitivity and specificity is tricky...
- Some further reading: Martinez et al 2008, Stærk-Østergaard et al 2022